

519095_Bosan_Individual2

2023-12-16

##Exercise 6.8: Problem 8 (parts e & f) ##e

```
library(glmnet)
```

```
## Warning: 套件 'glmnet' 是用 R 版本 4.3.2 來建造的
```

```
## 載入需要的套件 : Matrix
```


```
## Loaded glmnet 4.1-8
```

```
set.seed(15)
n <- 100
p <- 10
X <- matrix(rnorm(n * p), n, p)
y <- rnorm(n)

train_indices <- sample(1:n, size = 0.8 * n)
X_train <- X[train_indices, ]
y_train <- y[train_indices]
X_test <- X[-train_indices, ]
y_test <- y[-train_indices]

cv.out <- cv.glmnet(X_train, y_train, alpha = 1)

plot(cv.out)
```



```
best_lambda <- cv.out$lambda.min
lasso.mod <- glmnet(X_train, y_train, alpha = 1, lambda = best_lambda)
lasso.pred <- predict(lasso.mod, s = best_lambda, newx = X_test)

mse <- mean((lasso.pred - y_test)^2)
lasso.coef <- coef(lasso.mod, s = best_lambda)
print(lasso.coef)
```

```
## 11 x 1 sparse Matrix of class "dgCMatrix"
##              s1
## (Intercept) -0.07362781
## V1          0.00000000
## V2          .
## V3          .
## V4          .
## V5          .
## V6          .
## V7          .
## V8          .
## V9          .
## V10         .
```

```
print(mse)
```

```
## [1] 0.9363548
```

```
##Lazzo can remove some predictors
```

```
#f
```

```
library(glmnet)
set.seed(15)

beta_0 <- 1.5
beta_7 <- 2.0
Y <- beta_0 + beta_7 * X[, 7] + rnorm(n)

Y_train <- Y[train_indices]
Y_test <- Y[-train_indices]

lasso_mod_Y <- glmnet(X_train, Y_train, alpha = 1, lambda = best_lambda)
lasso_pred_Y <- predict(lasso_mod_Y, s = best_lambda, newx = X_test)

mse_Y <- mean((lasso_pred_Y - Y_test)^2)

lasso_coef_Y <- coef(lasso_mod_Y, s = best_lambda)
print(lasso_coef_Y)
```

```
## 11 x 1 sparse Matrix of class "dgCMatrix"
##              s1
## (Intercept) 1.5152170
## V1          0.8678231
## V2          .
## V3          .
## V4          .
## V5          .
## V6          .
## V7          1.8721727
## V8          .
## V9          .
## V10         .
```

```
print(mse_Y)
```

```
## [1] 0.05644347
```

```

best_subset <- function(X, Y, size) {
  n <- ncol(X)
  best_score <- Inf
  best_model <- NULL

  for (i in 1:size) {
    combinations <- combn(n, i, simplify = FALSE)
    for (comb in combinations) {
      X_subset <- as.matrix(X[, comb, drop = FALSE])
      # SKIP glmnet IF ONE COL
      if (ncol(X_subset) == 1) {
        next
      }
      model <- glmnet(X_subset, Y, alpha = 1, lambda = best_lambda)
      pred <- predict(model, newx = X_subset, s = best_lambda)
      mse <- mean((pred - Y)^2)

      if (mse < best_score) {
        best_score <- mse
        best_model <- comb
      }
    }
  }

  return(list("model" = best_model, "score" = best_score))
}

best_subset_result <- best_subset(X_train, Y_train, p)
print(best_subset_result)

```

```

## $model
## [1] 1 3 5 7
##
## $score
## [1] 0.03717781

```

The 1th, 3th, 5th, 7th predictors work best in this Lasso model, with only 0.0372 MSE

##Exercise 8.4: Problem 8 (parts a, b, & c) ##Problem #8: In the lab, a classification tree was applied to the Carseats data set after converting Sales into a qualitative ##response variable. Now we will seek to predict Sales using regression trees and related approaches, treating the response as ##a quantitative variable.

##(a) Split the data set into a training set and a test set.

```
library(ISLR)
```

```
## Warning: 套件 'ISLR' 是用 R 版本 4.3.2 來建造的
```

```
set.seed(1)

train = sample(1:nrow(Carseats), nrow(Carseats)/2)
car_train = Carseats[train, ]
car_test = Carseats[-train,]
```

##(b) Fit a regression tree to the training set. Plot the tree, and interpret the results. What test MSE do you obtain?

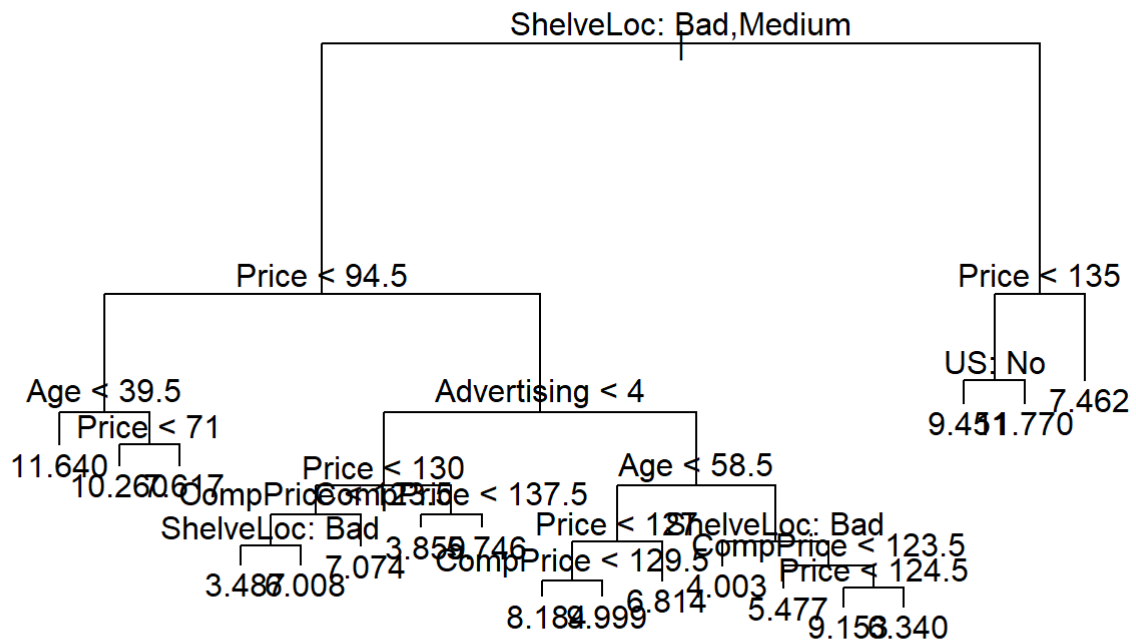
```
library(tree)
```

```
## Warning: 套件 'tree' 是用 R 版本 4.3.2 來建造的
```

```
# train the tree
tree_regression = tree(Sales~.,data = car_train)
summary(tree_regression)
```

```
##
## Regression tree:
## tree(formula = Sales ~ ., data = car_train)
## Variables actually used in tree construction:
## [1] "ShelveLoc" "Price" "Age" "Advertising" "CompPrice"
## [6] "US"
## Number of terminal nodes: 18
## Residual mean deviance: 2.167 = 394.3 / 182
## Distribution of residuals:
## Min. 1st Qu. Median Mean 3rd Qu. Max.
## -3.88200 -0.88200 -0.08712 0.00000 0.89590 4.09900
```

```
plot(tree_regression)
text(tree_regression ,pretty =0)
```



```

tree_prediction = predict(tree_regression, newdata=car_test)
tree_MSE <- mean((tree_prediction - car_test$Sales)^2)
tree_MSE

```


```
## [1] 4.922039
```

##(c) Use cross-validation in order to determine the optimal level of tree complexity. Does pruning the tree improve the test MSE?

```

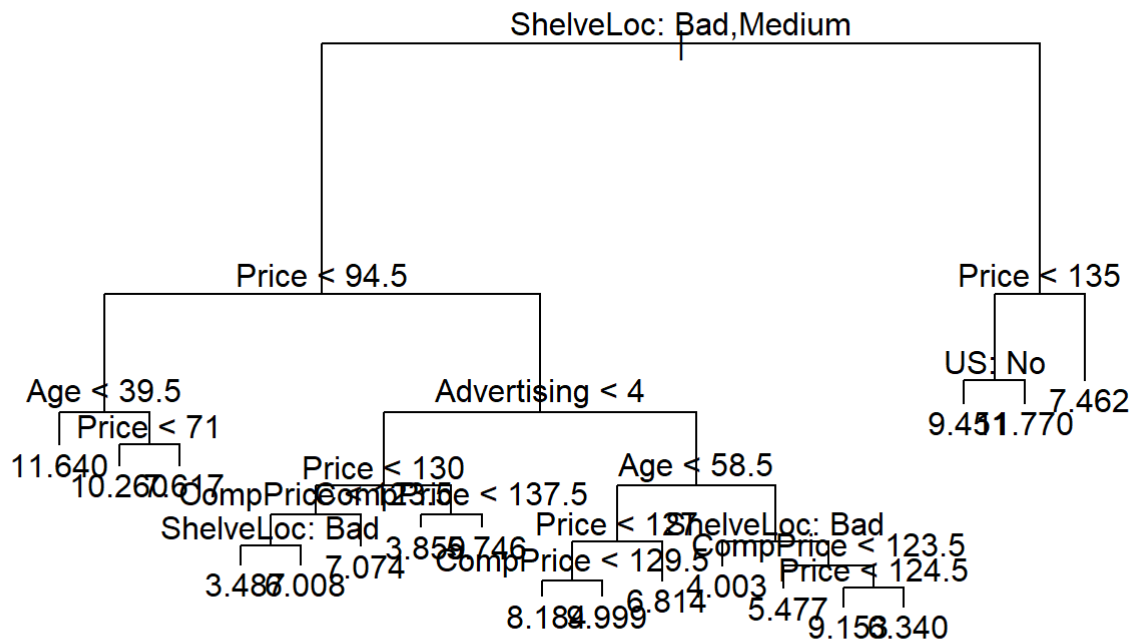
cv.carseats = cv.tree(tree_regression, FUN = prune.tree)
plot(cv.carseats$size, cv.carseats$dev, type = "b")

```



The graph above shows that deviation keep dropping when the size is 18 ## Thus, I picked 18 as the best size

```
prune_car = prune.tree(tree_regression, best = 18)
plot(prune_car)
text(prune_car, pretty = 0)
```



```
prune_prediction = predict(prune_car, newdata= car_test)

Prune_MSE = mean((prune_prediction - car_test$Sales)^2)
Prune_MSE
```

```
## [1] 4.922039
```

##The best size is 18. (Cross-Validation) ##In my test, the MSE remained the same after pruning the trees (remained 4.922039)

##Problem #8: In the lab, a classification tree was applied to the Carseats data set after converting Sales into a qualitative response variable. Now we will seek to predict Sales using regression trees and related approaches, treating the response as a quantitative variable. # (d) Use the bagging approach in order to analyze this data. What test MSE do you obtain? Use the importance() function to determine which variables are most important.

```
library(randomForest)
```

```
## Warning: 套件 'randomForest' 是用 R 版本 4.3.2 來建造的
```

```
## randomForest 4.7-1.1
```

```
## Type rfNews() to see new features/changes/bug fixes.
```



```
set.seed(88)
bagging_car = randomForest( Sales~., data = car_train, mtry = 10, importance = TRUE)
yhat_bagging = predict( bagging_car, newdata = car_test)
bagging_mse = mean((yhat_bagging - car_test$Sales)^2)
print(bagging_mse)
```

```
## [1] 2.618008
```

```
importance(bagging_car)
```

```
##           %IncMSE IncNodePurity
## CompPrice  27.909312    172.84191
## Income      6.098099     92.62640
## Advertising 13.203302    100.80551
## Population -2.264739     59.04411
## Price      56.594782    506.92540
## ShelfLoc   49.366388    369.15451
## Age        17.079423    157.21384
## Education   0.415839     44.67035
## Urban       1.710869     10.09492
## US          3.875395     17.49415
```

```
## The MSE using bagging is 2.618008
## Price's "%IncMSE" is 56.594782 and "it'sIncNodePurity" is 506.92540, which these both values are the highest values.
## Thus, Price is the most important variable in this case.
```

(e) Use random forests to analyze this data. What test MSE do you obtain? Use the importance() function to determine which variables are most important. Describe the effect of m, the number of variables considered at each split, on the error rate obtained.

```
set.seed(88)
randomforest_car = randomForest(Sales~., data = car_train, mtry = 3, importance = TRUE)
yhat_randomforest = predict(randomforest_car, newdata = car_test)
mse_randomforest = mean((yhat_randomforest - car_test$Sales)^2)
print(mse_randomforest)
```

```
## [1] 3.005177
```

```
importance(randomforest_car)
```

```
##           %IncMSE IncNodePurity
## CompPrice 13.7741065    151.14646
## Income    2.8594620    125.83963
## Advertising 8.4696416    108.00899
## Population -3.2377263    101.33362
## Price     35.7581303    389.53650
## ShelfLoc  34.7360147    293.78828
## Age       14.1990788    177.31774
## Education  2.0900556     76.26253
## Urban      0.3356185     16.59915
## US         5.1115996     33.41772
```

```
## The MSE using random forest is 3.005177 > 2.618008, there is no improvement using rf instead of bagging in this case
## Price's "%IncMSE" is 56.594782 and "it'sIncNodePurity" is 506.92540, which these both values are the highest values.
## Thus, Price is the most important variable in this case.
```

#Problem #10: We now use boosting to predict Salary in the Hitters data set. #(a) Remove the observations for whom the salary information is unknown, and then log-transform the salaries.

```
library(dplyr)
```

```
##
## 載入套件: 'dplyr'
```

```
## 下列物件被遮斷自 'package:randomForest':
##
##      combine
```

```
## 下列物件被遮斷自 'package:stats':
##
##      filter, lag
```

```
## 下列物件被遮斷自 'package:base':
##
##      intersect, setdiff, setequal, union
```

```
data("Hitters")

Hitters %>%
  filter(!is.na(Salary)) %>%
  mutate(Salary = log(Salary)) -> Hitters

head(Hitters)
```

##	AtBat	Hits	HmRun	Runs	RBI	Walks	Years	CAtBat	CHits	CHmRun
## -Alan Ashby	315	81	7	24	38	39	14	3449	835	69
## -Alvin Davis	479	130	18	66	72	76	3	1624	457	63
## -Andre Dawson	496	141	20	65	78	37	11	5628	1575	225
## -Andres Galarraga	321	87	10	39	42	30	2	396	101	12
## -Alfredo Griffin	594	169	4	74	51	35	11	4408	1133	19
## -Al Newman	185	37	1	23	8	21	2	214	42	1
##	CRuns	CRBI	CWalks	League	Division	PutOuts	Assists	Errors		
## -Alan Ashby	321	414	375	N	W	632	43	10		
## -Alvin Davis	224	266	263	A	W	880	82	14		
## -Andre Dawson	828	838	354	N	E	200	11	3		
## -Andres Galarraga	48	46	33	N	E	805	40	4		
## -Alfredo Griffin	501	336	194	A	W	282	421	25		
## -Al Newman	30	9	24	N	E	76	127	7		
##	Salary	NewLeague								
## -Alan Ashby	6.163315		N							
## -Alvin Davis	6.173786		A							
## -Andre Dawson	6.214608		N							
## -Andres Galarraga	4.516339		N							
## -Alfredo Griffin	6.620073		A							
## -Al Newman	4.248495		A							

- b. Create a training set consisting of the first 200 observations, and a test set consisting of the remaining observations.

```
hit_train <- Hitters[1:200,]
hit_test <- Hitters[-(1:200),]
```

- c. Perform boosting on the training set with 1,000 trees for a range of values of the shrinkage parameter λ . Produce a plot with different shrinkage values on the x-axis and the corresponding training set MSE on the y-axis.

```
library(gbm)
```

```
## Warning: 套件 'gbm' 是用 R 版本 4.3.2 來建造的
```

```
## Loaded gbm 2.1.8.1
```

```


shrinkage_values <- seq(0.01, 0.1, by = 0.01)

#store every model's error
train_mse <- rep(NA, length(shrinkage_values))

#boosting
for (i in seq_along(shrinkage_values)) {
  set.seed(10)
  gbm_model <- gbm(Salary ~ ., data = hit_train, distribution = "gaussian",
                    n.trees = 1000, shrinkage = shrinkage_values[i])
  pred <- predict(gbm_model, hit_train, n.trees=1000)
  train_mse[i] <- mean((pred - hit_train$Salary)^2)
}

plot(shrinkage_values, train_mse, type="b", xlab="Shrinkage", ylab="Training MSE")

```



- d. Produce a plot with different shrinkage values on the x-axis and the corresponding test set MSE on the y-axis.

```

library(gbm)


shrinkage_values <- seq(0.01, 0.1, by = 0.01)

test_mse <- rep(NA, length(shrinkage_values))

for (i in seq_along(shrinkage_values)) {
  set.seed(10)
  gbm_model <- gbm(Salary ~ ., data = hit_train, distribution = "gaussian",
                    n.trees = 1000, shrinkage = shrinkage_values[i])
  pred <- predict(gbm_model, hit_test, n.trees = 1000)
  test_mse[i] <- mean((pred - hit_test$Salary)^2)
}

plot(shrinkage_values, test_mse, type = "b", xlab = "Shrinkage", ylab = "Test MSE")

```



```

boosted_mse = min(test_mse)
boosted_mse

```

```
## [1] 0.2442988
```

- e. Compare the test MSE of boosting to the test MSE that results from applying two of the regression approaches seen in Chapters 3 and 6.

```
library(glmnet)

lm = lm(Salary ~ ., data = hit_train)
lm_prediction = predict(lm, newdata = hit_test)
lm_mse = mean((lm_prediction - hit_test$Salary)^2)

x_train <- model.matrix(Salary ~ ., hit_train)[,-1]
y_train <- hit_train$Salary
x_test <- model.matrix(Salary ~ ., hit_test)[,-1]
y_test <- hit_test$Salary

lasso <- glmnet(x_train, y_train, alpha = 1, lambda = 0.1)
lasso_predictions <- predict(lasso, s = 0.1, newx = x_test)
lasso_test_mse <- mean((lasso_predictions - y_test)^2)

print(paste("Linear Test MSE:", lm_mse))
```

```
## [1] "Linear Test MSE: 0.491795937545494"
```

```
print(paste("Lazzo Test MSE:", lasso_test_mse))
```

```
## [1] "Lazzo Test MSE: 0.43874517155745"
```


```
print(paste("Boosted Test MSE:", boosted_mse))
```

```
## [1] "Boosted Test MSE: 0.244298771115224"
```

```
#Boosted has the smallest MSE, 0.2443.
```

f. Which variables appear to be the most important predictors in the boosted model?

```
boosted.model = gbm(Salary~., data = hit_train, distribution = "gaussian", n.trees = 1000, shrinkage=shrinkage_values[which.min(test_mse)])
summary(boosted.model)
```



```
##          var    rel.inf
## CAtBat    CAtBat 18.4102380
## CRBI      CRBI  7.8057942
## Years     Years  7.7893112
## PutOuts   PutOuts 7.2161375
## Walks     Walks  7.1898528
## CHits     CHits  6.8821517
## CWalks    CWalks 6.4770790
## CHmRun    CHmRun 5.9336103
## Hits      Hits  5.5554983
## Assists   Assists 5.2785393
## CRuns     CRuns  4.8314369
## HmRun     HmRun  3.7926550
## AtBat     AtBat  3.5392703
## RBI       RBI   2.8738555
## Errors    Errors 2.5164422
## Runs      Runs  2.4744086
## Division  Division 0.6613331
## NewLeague NewLeague 0.6234775
## League    League 0.1489086
```

#CAtBat is the most important variable(highest rel.inf)

g. Now apply bagging to the training set. What is the test set MSE for this approach?

```
bagging = randomForest(Salary~., data = hit_train, distribution = "gaussian", n.trees = 500,
shrinkage = lambdas[which.min(test.error)], mtry = 19, importance = TRUE)
bagging_prediction = predict(bagging, hit_test)

bagging_test_mse = mean((bagging_prediction - hit_test$Salary)^2)
bagging_test_mse
```

```
## [1] 0.2298841
```

```
#The test MSE is 0.2321575 in this approach
```

#Proble, 11.4 #Direct Mailing to Airline Customers. East-West Airlines has entered into a partnership with the wireless phone company Telcon to sell the latter's service via direct mail. The file EastWestAirlinesNN.csv Download EastWestAirlinesNN.csv contains a subset of a data sample of who has already received a test offer. About 13% accepted.

#You are asked to develop a model to classify East-West customers as to whether they purchase a wireless phone service contract (outcome variable Phone_Sale). This model will be used to classify additional customers.

#1. Run a neural net model on these data, using a single hidden layer with 5 nodes. Remember to first convert categorical variables into dummies and scale numerical predictor variables to a 0-1 (use function preprocess() with method="range" - see Chapter 7). Generate a deciles-wise lift chart for the training and validation sets. Interpret the meaning (in business terms) of the leftmost bar of the validation decile- wise lift chart.

```
library(caret)
```

```
## Warning: 套件 'caret' 是用 R 版本 4.3.2 來建造的
```

```
## 載入需要的套件 : ggplot2
```

```
## Warning: 套件 'ggplot2' 是用 R 版本 4.3.2 來建造的
```

```
##
## 載入套件 : 'ggplot2'
```

```
## 下列物件被遮斷自 'package:randomForest':
##
##      margin
```

```
## 載入需要的套件 : lattice
```



```

library(nnet)
library(ggplot2)
library(dplyr)

# Load your data
data <- read.csv("EastWestAirlinesNN.csv")
data <- na.omit(data)

# Convert categorical variables into dummy variables
dummies <- dummyVars("~ .", data = data)
data_transformed <- predict(dummies, newdata = data)

# Scale numerical predictor variables to a 0-1 range
preproc <- preProcess(data_transformed, method = "range")
data_scaled <- predict(preproc, data_transformed)

if (!is.data.frame(data_scaled)) {
  data_scaled <- as.data.frame(data_scaled)
}

if (!("Phone_sale" %in% names(data_scaled))) {
  data_scaled$Phone_sale <- data$Phone_sale
}

# Split the data into training and validation sets
set.seed(123) # for reproducibility
trainingIndex <- createDataPartition(data_scaled$Phone_sale, p = .8, list = TRUE)
trainingData <- data_scaled[trainingIndex[[1]], ]
validationData <- data_scaled[-trainingIndex[[1]], ]

```

```

# Load necessary libraries
library(caret)
library(nnet)
library(ggplot2)
library(dplyr)
#install.packages("neuralnet")
library(neuralnet)

```

```
## Warning: 套件 'neuralnet' 是用 R 版本 4.3.2 來建造的
```

```
##
## 載入套件: 'neuralnet'
```

```
## 下列物件被遮斷自 'package:dplyr':
##
##      compute
```

```

# Load your data
data <- read.csv("EastWestAirlinesNN.csv")
data <- na.omit(data)

# Convert categorical variables into dummy variables
dummies <- dummyVars("~ .", data = data)
data_transformed <- predict(dummies, newdata = data)
data_transformed <- as.data.frame(data_transformed)
data_transformed$Phone_sale <- as.factor(data$Phone_sale)

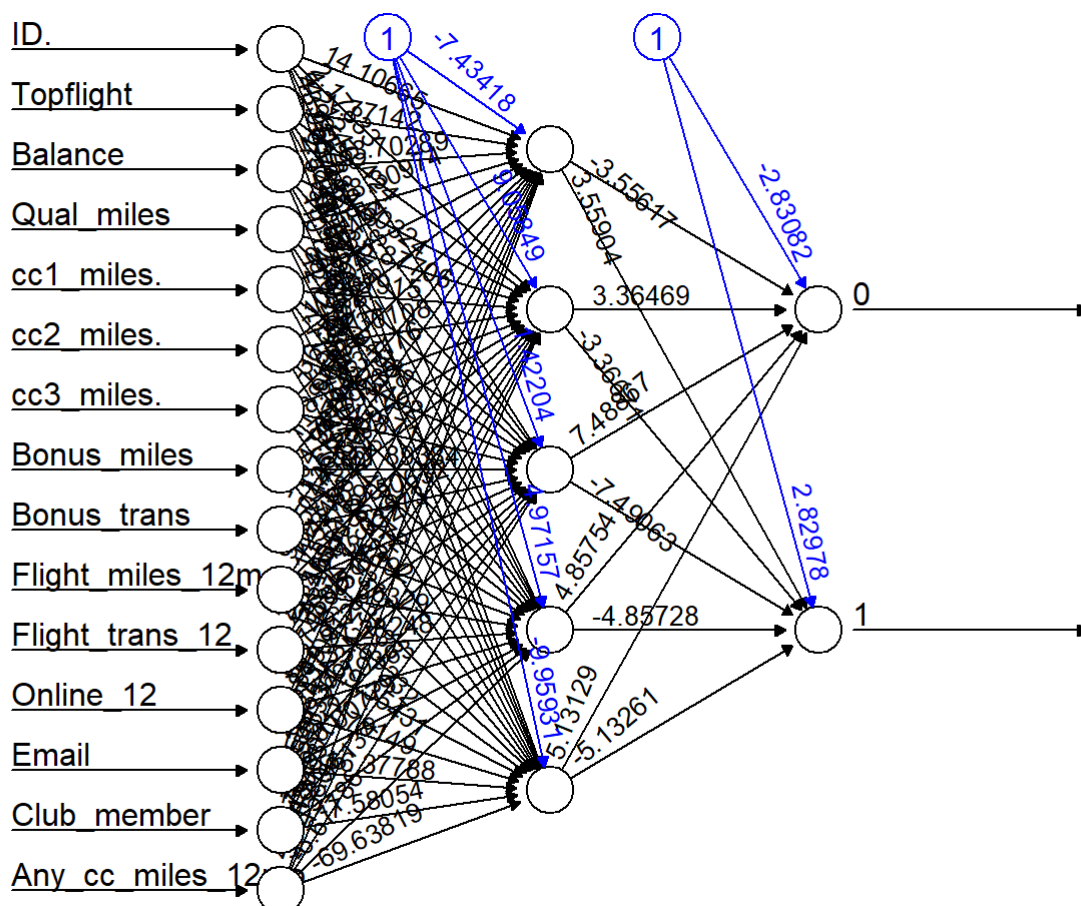
#scale 0-1
preproc <- preProcess(data_transformed[, -which(names(data_transformed) == "Phone_sale")], method = "range")
data_scaled <- predict(preproc, data_transformed)

data_scaled <- as.data.frame(data_scaled)
data_scaled$Phone_sale <- as.factor(data$Phone_sale)

# Split the data into training and validation sets
set.seed(123) # for reproducibility
trainingIndex <- createDataPartition(data_scaled$Phone_sale, p = .8, list = TRUE)
trainingData <- data_scaled[trainingIndex[[1]], ]
validationData <- data_scaled[-trainingIndex[[1]], ]


nn_model <- neuralnet(Phone_sale ~ ., data = trainingData, hidden = 5, linear.output = FALSE)
plot(nn_model, rep = "best")

```



```
#install.packages("gains")
library(gains)
prediction <- predict(nn_model, trainingData, type = "raw")
actual_numeric <- as.numeric(as.character(trainingData$Phone_sale))
predicted_probabilities <- prediction[, 2]
gain <- gains(actual_numeric, predicted_probabilities)
barplot(gain$mean.resp / mean(actual_numeric), names.arg = gain$depth,
        xlab = "%", ylab = "mean", main = "Training Lift Chart")
```

Training Lift Chart




The first bar in a decile-wise lift chart represents the top 10% of cases predicted by a model to be the most likely to have the positive outcome you're interested in.
 ## For example, this could represent a group of customers most likely to respond to a marketing campaign, the most profitable segment, or those who are most likely to churn, depending on the context of the model's objective.

##2. Comment on the difference between the training and validation lift charts.

```
prediction_valid <- predict(nn_model, validationData, type = "raw")
actual_numeric_valid <- as.numeric(as.character(validationData$Phone_sale))
predicted_probabilities_valid <- prediction_valid[, 2]
gain_valid <- gains(actual_numeric_valid, predicted_probabilities_valid)
barplot(gain_valid$mean.resp / mean(actual_numeric_valid), names.arg = gain_valid$depth,
        xlab = "%", ylab = "mean", main = "Validation Lift Chart")
```

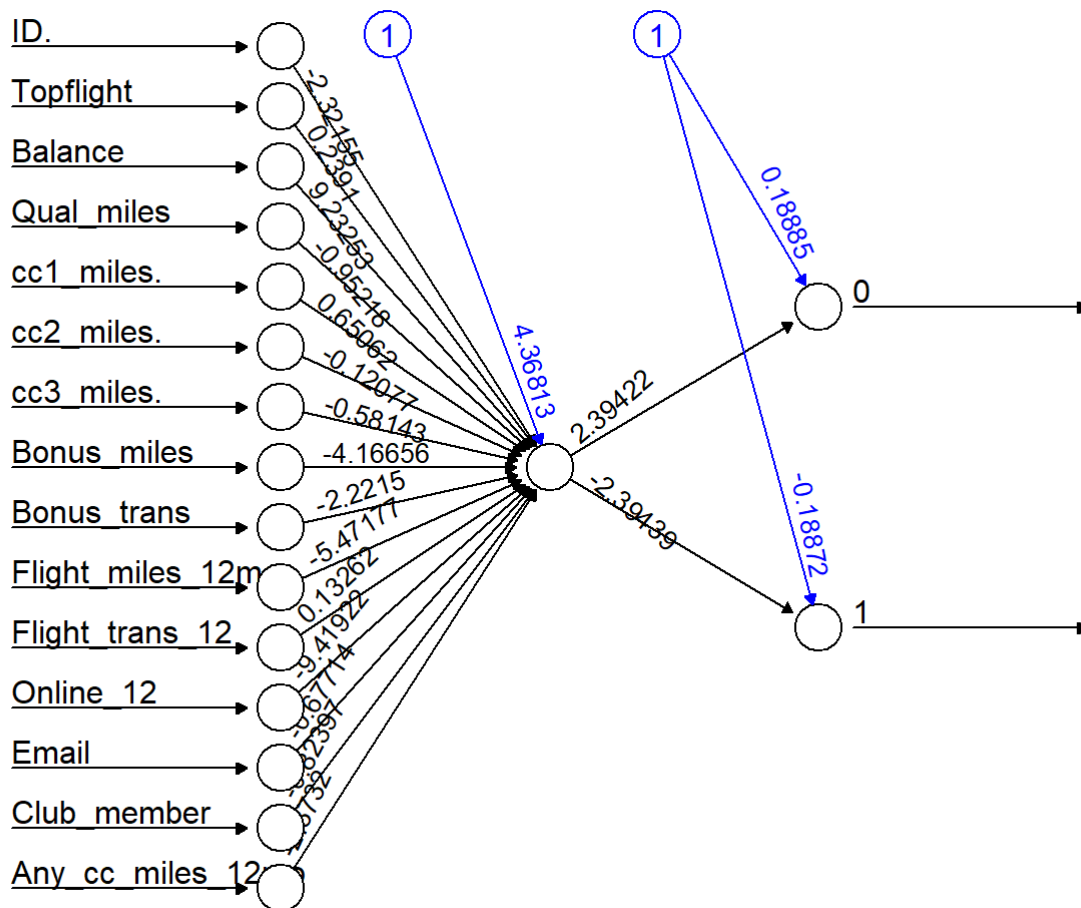
Validation Lift Chart



The training lift chart shows a significantly higher lift in the first decile compared to the validation lift chart. This suggests that the model may overfits to the training data.

3. Run a second neural net model on the data, this time setting the number of hidden nodes to 1. Comment now on the difference between this model and the model you ran earlier, and how overfitting might have affected results.

```
nn_model_2 <- neuralnet(Phone_sale ~ ., data = trainingData, hidden = 1, linear.output = FALSE)
plot( nn_model_2, rep = "best")
```



```
error_nn_model <- as.character(nn_model$result.matrix[1,])
print(paste("hidden = 5, error =", error_nn_model))
```

```
## [1] "hidden = 5, error = 400.95224779195"
```

```
error_nn_model_2 <- as.character(nn_model_2$result.matrix[1,])
print(paste("hidden = 1, error =", error_nn_model_2))
```

```
## [1] "hidden = 1, error = 436.320567915564"
```

```
## when there is 5 hidden nodes, the error is lower
## however, having more node might overfit the data
```

4. What sort of information, if any, is provided about the effects of the various variables?

```
nn_model$result.matrix
```

```

##                                     [,1]
## error                             4.009522e+02
## reached.threshold                 9.859623e-03
## steps                             1.412400e+04
## Intercept.to.1layhid1             -7.434181e+00
## ID..to.1layhid1                   1.410665e+01
## Topflight.to.1layhid1              1.771420e+00
## Balance.to.1layhid1                1.137029e+02
## Qual_miles.to.1layhid1             -3.783097e+02
## cc1_miles..to.1layhid1             4.649399e-01
## cc2_miles..to.1layhid1            -1.009413e+01
## cc3_miles..to.1layhid1            -3.629226e+01
## Bonus_miles.to.1layhid1            1.189053e+02
## Bonus_trans.to.1layhid1           -1.049871e+02
## Flight_miles_12mo.to.1layhid1      9.350446e+01
## Flight_trans_12.to.1layhid1        2.241153e+01
## Online_12.to.1layhid1              3.511434e+02
## Email.to.1layhid1                 2.258518e+00
## Club_member.to.1layhid1           -7.258084e+01
## Any_cc_miles_12mo.to.1layhid1      4.357676e+01
## Intercept.to.1layhid2             9.058490e+00
## ID..to.1layhid2                   2.178335e+00
## Topflight.to.1layhid2             -5.704240e+00
## Balance.to.1layhid2                3.646032e+02
## Qual_miles.to.1layhid2            -4.358771e+02
## cc1_miles..to.1layhid2             8.219150e+01
## cc2_miles..to.1layhid2             1.330108e+01
## cc3_miles..to.1layhid2            -2.333176e+01
## Bonus_miles.to.1layhid2           -7.279435e+01
## Bonus_trans.to.1layhid2           -1.233009e+02
## Flight_miles_12mo.to.1layhid2      1.244364e+02
## Flight_trans_12.to.1layhid2        2.340551e+01
## Online_12.to.1layhid2              6.774487e+01
## Email.to.1layhid2                 -1.073709e+01
## Club_member.to.1layhid2           -7.854633e+01
## Any_cc_miles_12mo.to.1layhid2     -1.748063e+01
## Intercept.to.1layhid3             1.422043e+00
## ID..to.1layhid3                   4.212456e+00
## Topflight.to.1layhid3              6.328951e+01
## Balance.to.1layhid3                3.473118e+01
## Qual_miles.to.1layhid3            -2.376700e+01
## cc1_miles..to.1layhid3            -3.091388e+01
## cc2_miles..to.1layhid3             5.405492e+01
## cc3_miles..to.1layhid3             5.418715e+00
## Bonus_miles.to.1layhid3           -5.118938e+02
## Bonus_trans.to.1layhid3           -1.995063e+02
## Flight_miles_12mo.to.1layhid3     -8.970327e+02
## Flight_trans_12.to.1layhid3        3.593113e+02
## Online_12.to.1layhid3             -8.991200e+01
## Email.to.1layhid3                 -4.136631e+01
## Club_member.to.1layhid3            4.390112e+01
## Any_cc_miles_12mo.to.1layhid3      3.845806e+01
## Intercept.to.1layhid4             4.971570e+00
## ID..to.1layhid4                   -4.991784e+00
## Topflight.to.1layhid4              2.582949e+00

```

```

## Balance.to.1layhid4      1.397843e+01
## Qual_miles.to.1layhid4  -9.053788e+00
## cc1_miles..to.1layhid4   7.386216e+00
## cc2_miles..to.1layhid4  -1.700824e+00
## cc3_miles..to.1layhid4   4.081439e+01
## Bonus_miles.to.1layhid4  -3.269719e+00
## Bonus_trans.to.1layhid4   4.341923e+00
## Flight_miles_12mo.to.1layhid4 -3.632895e-01
## Flight_trans_12.to.1layhid4  1.128248e+01
## Online_12.to.1layhid4    -1.193627e+00
## Email.to.1layhid4        -7.786017e-03
## Club_member.to.1layhid4   -2.961337e+00
## Any_cc_miles_12mo.to.1layhid4 -6.677862e+00
## Intercept.to.1layhid5    -9.959312e+00
## ID..to.1layhid5          4.395670e+01
## Topflight.to.1layhid5    -7.549098e+02
## Balance.to.1layhid5      -1.012107e+03
## Qual_miles.to.1layhid5    1.272584e+03
## cc1_miles..to.1layhid5    -1.167898e+03
## cc2_miles..to.1layhid5    -3.125285e+00
## cc3_miles..to.1layhid5    -1.021083e+03
## Bonus_miles.to.1layhid5    1.373793e+03
## Bonus_trans.to.1layhid5   -5.466323e+02
## Flight_miles_12mo.to.1layhid5 -6.726123e+02
## Flight_trans_12.to.1layhid5  2.393543e+02
## Online_12.to.1layhid5     7.791486e+00
## Email.to.1layhid5         1.637788e+01
## Club_member.to.1layhid5    -7.580537e+00
## Any_cc_miles_12mo.to.1layhid5 -6.963819e+01
## Intercept.to.0           -2.830816e+00
## 1layhid1.to.0            -3.556167e+00
## 1layhid2.to.0            3.364686e+00
## 1layhid3.to.0            7.488670e+00
## 1layhid4.to.0            4.857535e+00
## 1layhid5.to.0            5.131289e+00
## Intercept.to.1           2.829779e+00
## 1layhid1.to.1            3.559039e+00
## 1layhid2.to.1            -3.366714e+00
## 1layhid3.to.1            -7.490632e+00
## 1layhid4.to.1            -4.857282e+00
## 1layhid5.to.1            -5.132612e+00

```

#The variance generalized weights are provided

#If the absolute value of the weight is high, it means that the variable has great impact on the outcome


Exercise 10.7: Problem 9

ISLR p.417

9. Consider the USArrests data. We will now perform hierarchical clustering on the states.

a. Using hierarchical clustering with complete linkage and Euclidean distance, cluster the states.

```
data("USArrests")
cluster_complete <- hclust(dist(USArrests), method = "complete")
plot(cluster_complete, main = "Hierarchical Clustering with Complete Linkage", sub = "", xlab = "")
```



b. Cut the dendrogram at a height that results in three distinct clusters. Which states belong to which clusters?

```
# From my perspective, I'll make split the data into three groups
cut_complete <- cutree(cluster_complete, k = 3)

clusters <- data.frame(State = names(cut_complete), Cluster = cut_complete)
clusters_sorted <- clusters[order(clusters$Cluster, clusters$State), ]
print(clusters_sorted)
```


##	State	Cluster
##	Alabama	1
##	Alaska	1
##	Arizona	1
##	California	1
##	Delaware	1
##	Florida	1
##	Illinois	1
##	Louisiana	1
##	Maryland	1
##	Michigan	1
##	Mississippi	1
##	Nevada	1
##	New Mexico	1
##	New York	1
##	North Carolina	1
##	South Carolina	1
##	Arkansas	2
##	Colorado	2
##	Georgia	2
##	Massachusetts	2
##	Missouri	2
##	New Jersey	2
##	Oklahoma	2
##	Oregon	2
##	Rhode Island	2
##	Tennessee	2
##	Texas	2
##	Virginia	2
##	Washington	2
##	Wyoming	2
##	Connecticut	3
##	Hawaii	3
##	Idaho	3
##	Indiana	3
##	Iowa	3
##	Kansas	3
##	Kentucky	3
##	Maine	3
##	Minnesota	3
##	Montana	3
##	Nebraska	3
##	New Hampshire	3
##	North Dakota	3
##	Ohio	3
##	Pennsylvania	3
##	South Dakota	3
##	Utah	3
##	Vermont	3
##	West Virginia	3
##	Wisconsin	3

```
#Alabama, Alaska, Arizona, California, Delaware, Florida, Illinois, Louisiana, Maryland, Michigan, Mississippi, Nevada, New Mexico, New York, Carolina North Carolina: Cluster 1
```

```
# Arkansas, Colorado, Georgia, Massachusetts, Missouri, New Jersey, Oklahoma, Oregon, Tennessee, Texas, Virginia, Washington Washington 2
```

```
#Others are in the third cluster
```

- c. Hierarchically cluster the states using complete linkage and Euclidean distance, after scaling the variables to have standard deviation one.

```
library(purrr)
```

```
## Warning: 套件 'purrr' 是用 R 版本 4.3.2 來建造的
```

```
##  
## 載入套件: 'purrr'
```

```
## 下列物件被遮斷自 'package:caret':  
##  
## lift
```

```
scaling <- function(x) (x - mean(x, na.rm = TRUE)) / sd(x, na.rm = TRUE)
```

```
scaled_cluster <- USArrests %>%  
  map_df(scaling) %>%  
  dist(method = 'euclidean') %>%  
  hclust(method = 'complete')
```

```
scaled_cluster$labels <- row.names(USArrests)[scaled_cluster$order]  
scaled_cluster$labels
```

```
## [1] "South Dakota" "West Virginia" "North Dakota" "Vermont"  
## [5] "Maine" "Iowa" "New Hampshire" "Idaho"  
## [9] "Montana" "Nebraska" "Kentucky" "Arkansas"  
## [13] "Virginia" "Wyoming" "Missouri" "Oregon"  
## [17] "Washington" "Delaware" "Rhode Island" "Massachusetts"  
## [21] "New Jersey" "Connecticut" "Minnesota" "Wisconsin"  
## [25] "Oklahoma" "Indiana" "Kansas" "Ohio"  
## [29] "Pennsylvania" "Hawaii" "Utah" "Colorado"  
## [33] "California" "Nevada" "Florida" "Texas"  
## [37] "Illinois" "New York" "Arizona" "Michigan"  
## [41] "Maryland" "New Mexico" "Alaska" "Alabama"  
## [45] "Louisiana" "Georgia" "Tennessee" "North Carolina"  
## [49] "Mississippi" "South Carolina"
```

- d. What effect does scaling the variables have on the hierarchical clustering obtained? In your opinion, should the variables be scaled before the inter-observation dissimilarities are computed? Provide a justification for your answer.

Scaling the variables ensures that each variable contributes equally to the distance calculations

Without scaling, there might be dominating variables, and the clustering results might be skewed towards variables with larger magnitudes.

We should scale the variables before we compute the euclidean distances

So that there will be no dominating variables, and the clustering process will be more interpretable, as it removes the bias introduced by the scale of the variables.