

## Exercise 8.4

Section 23 - Group 6 Project Groups (Bosan Hsu, Fan Liu, Jimeng Yin, Michael Liu, Richard Wang, Zhuoqian Zhang)

```
library(ISLR)
attach(Default)
set.seed(1)
names(Default)

## [1] "default" "student" "balance" "income"

train = sample(1:nrow(Default), nrow(Default)/2)
Default.train = Default[train, ]
Default.test = Default[-train, ]

glm.Default = glm(default~., data=Default, subset=train, family = binomial)
glm.probs = predict(glm.Default, newdata = Default.test, type = "response")
glm.pred = ifelse(glm.probs > 0.5, 'Yes', 'No')
table(Default.test$default, glm.pred)

##      glm.pred
##      No  Yes
## No  4825  18
## Yes   112  45

mean(Default.test$default != glm.pred)

## [1] 0.026

library(randomForest)

## Warning: 程辑包 'randomForest' 是用 R 版本 4.3.2 来建造的

## randomForest 4.7-1.1

## Type rfNews() to see new features/changes/bug fixes.

set.seed(1)
bag.Default = randomForest(default~., data=Default, subset=train, mtry=3, ntree=250)
bag.probs = predict(bag.Default, newdata = Default.test, type = "prob")
bag.pred = ifelse(bag.probs[,2] > 0.5, 'Yes', 'No')
table(Default.test$default, bag.pred)

##      bag.pred
##      No  Yes
```

```

##      No  4799   44
##      Yes  108   49

mean(Default.test$default != bag.pred)

## [1] 0.0304

set.seed(1)
rf.Default = randomForest(default~., data=Default, subset=train, mtry=2,
  ntree=250)
rf.probs = predict(rf.Default, newdata = Default.test, type = "prob")
rf.pred = ifelse(rf.probs[,2] > 0.5, 'Yes', 'No')
table(Default.test$default, rf.pred)

##      rf.pred
##           No  Yes
##      No  4805   38
##      Yes  108   49

mean(Default.test$default != rf.pred)

## [1] 0.0292

library(gbm)

## Warning: 程辑包 'gbm'是用 R 版本 4.3.2 来建造的

## Loaded gbm 2.1.8.1

Default.train$default = as.numeric(Default.train$default) - 1
Default.test$default = as.numeric(Default.test$default) - 1
set.seed(1)
boost.fit = gbm(default~., data = Default.train, distribution = "bernoulli", n.trees = 250)
boost.probs = predict(boost.fit, newdata = Default.test, type = "response")

## Using 250 trees...

boost.pred = ifelse(boost.probs > 0.5, 1, 0)
table(Default.test$default, boost.pred)

##      boost.pred
##           0    1
##      0  4816   27
##      1   106   51

mean(Default.test$default != boost.pred)

## [1] 0.0266

set.seed(1)
lm.Default <- lm(default ~ ., data = Default.train)

```

```

lm.probs <- predict(lm.Default, newdata = Default.test, type = "response")
lm.pred <- ifelse(lm.probs > 0.5, 1, 0)
table(Default.test$default, lm.pred)

##      lm.pred
##           0
##    0 4843
##    1  157

mean(Default.test$default != lm.pred)

## [1] 0.0314

```

The accuracy of logistic is the best. Linear is the worst.